

# Report for Sep 29

## Introduction

This report focuses on testing a dataset with  $D = 1$  using the Stochastic Gradient Descent (SGD) method. In the current phase, we modified the forward function by introducing a new `nn.Module`, simulating `DivideByConstantLayer` instead of directly dividing by  $m$  after applying  $\langle \vec{a}, \cdot \rangle$ .

```
import torch
import torch.nn as nn

class DivideByConstantLayer(nn.Module):
    def __init__(self, constant):
        super(DivideByConstantLayer, self).__init__()
        self.constant = constant

    def forward(self, x):
        return x / self.constant
```

The model is defined as follows, employing HE initialization:

```
class Model(nn.Module):
    def __init__(self, m, D):
        super(Model, self).__init__()
        self.linear1 = torch.nn.Linear(D + 1, m)
        self.tanh = torch.nn.Tanh()
        self.linear2 = torch.nn.Linear(m, 1)
        self.divide = DivideByConstantLayer(m)

        # HE initialization
        torch.nn.init.kaiming_normal_(self.linear1.weight)
        torch.nn.init.kaiming_normal_(self.linear2.weight)

    def forward(self, x):
        fc1 = self.linear1(x)
        fc2 = self.tanh(fc1)
        fc3 = self.linear2(fc2)
        fc4 = self.divide(fc3)
        return fc4
```

The parameters are set as follows:  $D = 1$ ,  $m = 50$ , batch size = 32.

The relative error is computed as:

$$\text{error} = \frac{|y_{\text{pred}} - y_{\text{true}}|}{|y_{\text{true}}|}$$

## Stochastic Gradient Descent (SGD) Method

We selected PyTorch's SGD optimizer for training.

## Default Learning Rate

Initially, we tested the default learning rate with 100,000 epochs:

```
optimizer = torch.optim.SGD(model.parameters(), lr=learning_rate)
```

After 10,000 epochs, the results were:

- Train Loss: 0.005463460925966501
- Test Loss: 0.002377670258283615
- Relative Error: 0.15652328729629517

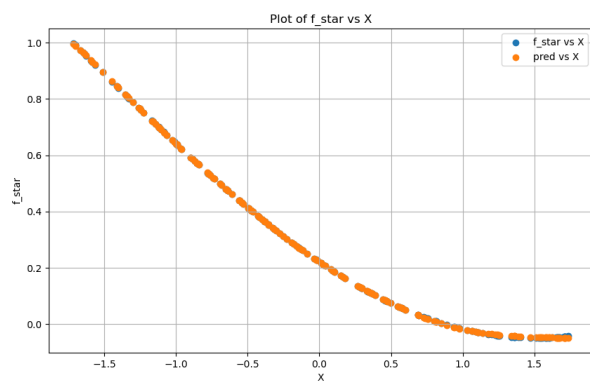
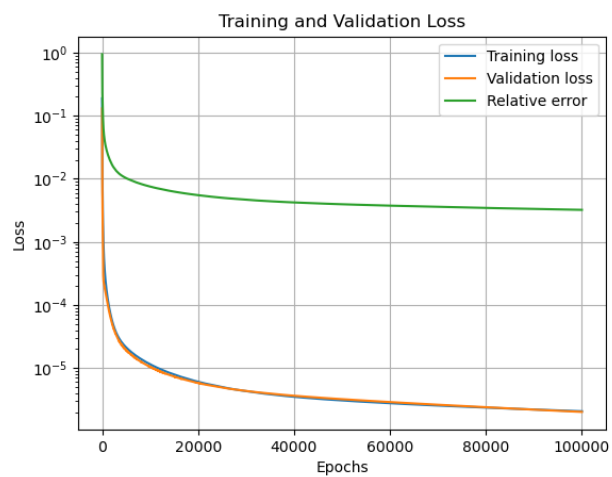
The default SGD method failed to adequately fit the curve. We then adjusted the learning rate to find an optimal result.

## Learning Rate = 1

After 100,000 epochs:

- Train Loss:  $2.066 \times 10^{-6}$
- Test Loss:  $2.018 \times 10^{-6}$
- Relative Error: 0.003224582178518176

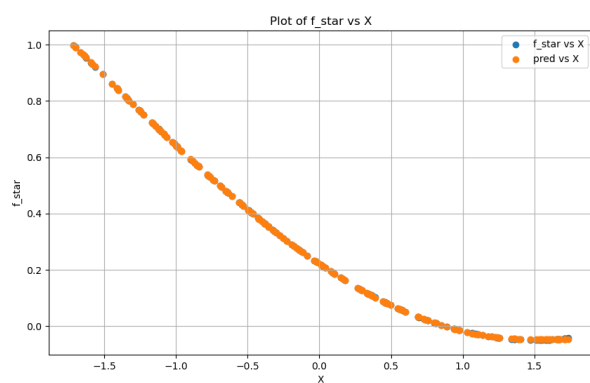
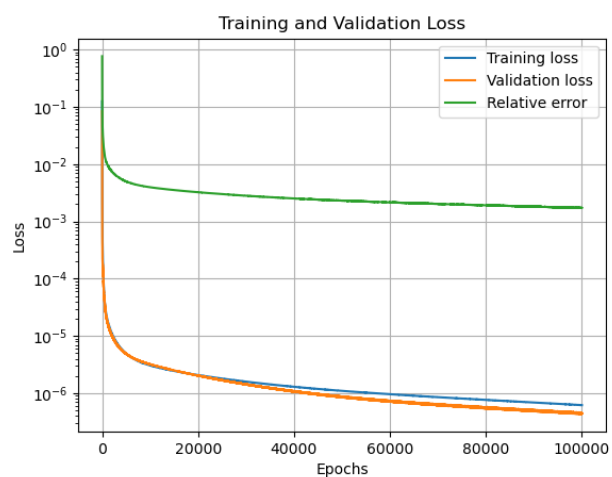
While increasing the learning rate improved the results, we observed significant errors in the interval  $x \in [1, 1.5]$ .



## Learning Rate = 5

After 100,000 epochs:

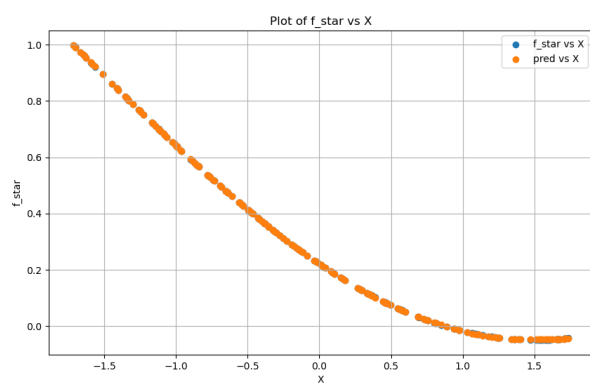
- Train Loss:  $6.208 \times 10^{-7}$
- Test Loss:  $4.469 \times 10^{-7}$
- Relative Error: 0.0017211942467838526



## Learning Rate = 10

After 100,000 epochs:

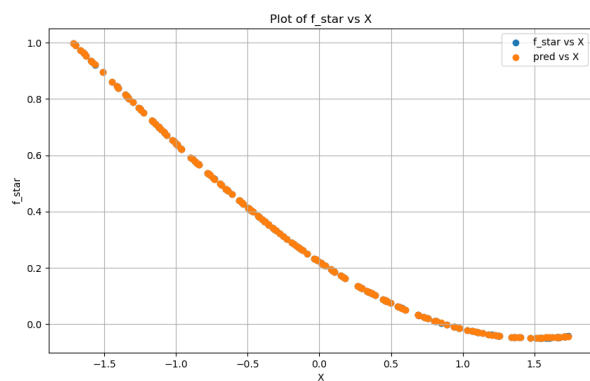
- Train Loss:  $2.060 \times 10^{-7}$
- Test Loss:  $1.590 \times 10^{-7}$
- Relative Error: 0.0009969599777832627



## Learning Rate = 15

After 100,000 epochs:

- Train Loss:  $9.079 \times 10^{-8}$
- Test Loss:  $7.902 \times 10^{-8}$
- Relative Error: 0.0006686306442134082

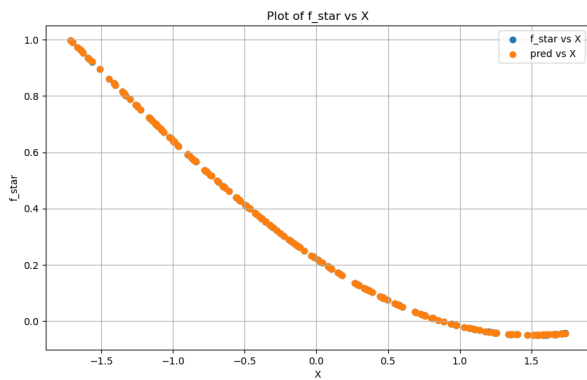


## Learning Rate = 20

After 100,000 epochs:

- Train Loss:  $5.602 \times 10^{-8}$
- Test Loss:  $5.357 \times 10^{-8}$
- Relative Error: 0.0005298249307088554

Significant oscillations were observed in both training and testing losses.

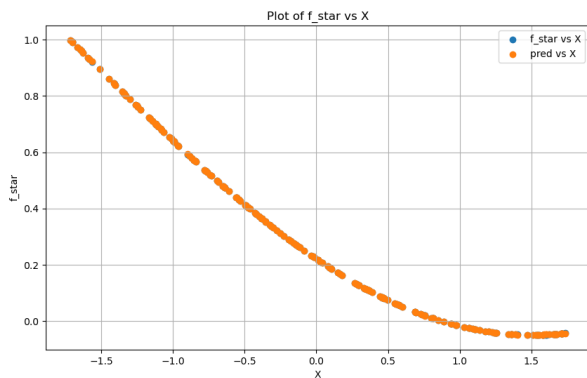
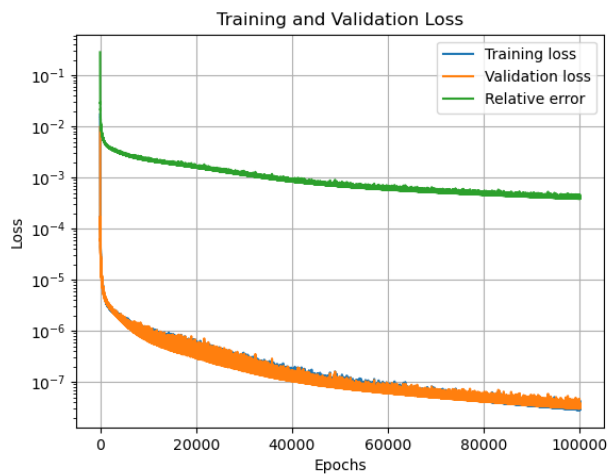


## Learning Rate = 30

After 100,000 epochs:

- Train Loss:  $2.965 \times 10^{-8}$
- Test Loss:  $3.113 \times 10^{-8}$
- Relative Error: 0.00038910453440621495

Further increasing the learning rate showed diminishing returns, with the relative error converging to approximately  $3 \times 10^{-4}$ . At this stage, we proceeded to experiment with the Adam optimizer.



## Adam Method

```
optimizer = torch.optim.Adam(model.parameters())
```

Using the default Adam optimizer, we observed the following at Epoch 65,574:

- Train Loss:  $1.883 \times 10^{-9}$
- Test Loss:  $1.899 \times 10^{-9}$
- Relative Error:  $9.765 \times 10^{-5}$

Given that the relative error dropped below  $1 \times 10^{-4}$ , we terminated the process.



